

M2-L1 Problem 1 (5 points)

In this question you will perform linear least squares regression on a very small dataset of 3 points. First, load and plot the data by running the following cell.


The variables provided are:

- x : 3x1 input data
- y : 3x1 output data

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

x = np.array([-1, 3, 5]).T
y = np.array([2, -1, 4]).T

fig, ax = plt.subplots()
plt.plot(x, y, 'o')
ax.spines['left'].set_position(('data', 0))
ax.spines['bottom'].set_position(('data', 0))
sns.despine()
plt.grid()
plt.show()
```



Construct a design matrix

For 1-D linear regression, the design matrix must contain not only a column of input x -values, but also a 'bias column' -- a column of ones (to allow the regression line to have an intercept).

The next step is to construct the design matrix X by concatenating a column of ones to the given input x . This has been done for you below:

```
In [2]: bias = np.ones_like(x)

X = np.concatenate([x,bias],1)

print("Design Matrix:\n",X)
```

Design Matrix:

```
[[ -1  1]
 [ 3  1]
 [ 5  1]]
```

Solving for regression coefficients

Now that we have the design matrix X and the output y , we can solve for the coefficients w such that $Xw \approx y$ using:

$$w = (X'X)^{-1}X'y$$

Note that you can use the following in Python:

- `@` for matrix multiplication
- `np.linalg.inv(A)` for inversion of matrix A
- `A.T` for transpose of a matrix A
- `b.reshape(-1,1)` to treat 1D array b as a column (you will need to do this for y)

Your line's slope should be ≈ 0.18 and your y-intercept should be ≈ 1.25 .

```
In [3]: # YOUR CODE GOES HERE
# Get coefficients w
w = np.linalg.inv(X.T@X)@X.T@y
print("Linear Coefficients:\n", w)
```

Linear Coefficients:

```
[[0.17857143]
 [1.25      ]]
```

Making predictions

Now that we have the coefficients, we can make predictions on new data with the model.

Do the following steps:

- [Given] Sample 40 points on the interval [-3,7], such as by using `np.linspace()` (Append `.reshape(-1,1)` to convert to a column)
- [Given] Create a design matrix by adding a column of ones as done previously
- Make a prediction by multiplying your new design matrix by `w`. You can do matrix multiplication with the `@` symbol
- [Given] Add a line to the plot showing these predictions

```
In [4]: n = 40
x_test = np.linspace(-4,7,n).reshape(-1,1)
bias_test = np.ones_like(x_test)
X_test = np.concatenate([x_test, bias_test], 1)

# YOUR CODE GOES HERE
# Predict y_test

y_test = X_test@w

fig, ax = plt.subplots()
plt.plot(x, y, '.')
plt.plot(x_test, y_test)
ax.spines['left'].set_position(('data', 0))
ax.spines['bottom'].set_position(('data', 0))
sns.despine()
plt.grid()
plt.show()
```

