

M3-L1 Problem 3 (6 points)

```
In [1]: import numpy as np
import matplotlib.pyplot as plt

def plot_data(data, c, title="", xlabel="$x_1$", ylabel="$x_2$", classes=["", ""], alpha=0.5):
    N = len(c)
    colors = ['royalblue', 'crimson']
    symbols = ['o', 's']

    plt.figure(figsize=(5, 5), dpi=120)

    for i in range(2):
        x = data[:, 0][c==i]
        y = data[:, 1][c==i]

        plt.scatter(x, y, color=colors[i], marker=symbols[i], edgecolor="black", linewidth=1, alpha=alpha)

    plt.legend(loc="upper right")
    plt.xlabel(xlabel)
    plt.ylabel(ylabel)
    ax = plt.gca()
    ax.set_xticklabels([])
    ax.set_yticklabels([])
    plt.xlim([-0.05, 1.05])
    plt.ylim([-0.05, 1.05])
    plt.title(title)

def plot_contour(predict, mapXY = None):
    res = 500
    vals = np.linspace(-0.05, 1.05, res)
    x, y = np.meshgrid(vals, vals)
    XY = np.concatenate((x.reshape(-1, 1), y.reshape(-1, 1)), axis=1)
    if mapXY is not None:
        XY = mapXY(XY)
    contour = predict(XY).reshape(res, res)
    plt.contour(x, y, contour)
```


Generate Dataset

(Don't edit this code.)

```
In [2]: def sample_ring(N, x, y, ro, ri):
        theta = np.random.rand(N)*2*np.pi
        r = np.random.rand(N)
        r = np.sqrt(r*(ro**2-ri**2)+ri**2)
        xs = x + r * np.cos(theta)
        ys = y + r * np.sin(theta)
        return xs, ys

def get_ring_dataset():
    np.random.seed(0)
    c0 = sample_ring(70, 0.5, 0.5, 0.5, 0.3)
    c1 = sample_ring(60, 0.45, 0.47, 0.36, 0.15)
    xs = np.concatenate([c0[0], c1[0]], 0)
    ys = np.concatenate([c0[1], c1[1]], 0)
    c = np.concatenate([np.zeros(70), np.ones(60)], 0)
    return np.vstack([xs, ys]).T, c
```

```
In [3]: data, classes = get_ring_dataset()
format = dict(xlabel="Initial Speed", ylabel="Launch Angle", classes=["0 - Misses ta
plot_data(data, classes, **format)
```



Feature Expansion

Define a function to expand 2 features into more features For the features x_1 and x_2 , expand into:

- 1
- x_1
- x_2
- x_1^2
- x_2^2
- $\sin(x_1)$
- $\cos(x_1)$
- $\sin(x_2)$
- $\cos(x_2)$
- $\sin^2(x_1)$
- $\cos^2(x_1)$
- $\sin^2(x_2)$
- $\cos^2(x_2)$
- $\exp(x_1)$
- $\exp(x_2)$

```
In [4]: def feature_expand(x):
        x1 = x[:,0].reshape(-1, 1)
        x2 = x[:,1].reshape(-1, 1)

        # YOUR CODE GOES HERE:
        columns = [np.ones_like(x1), x1, x2, x1**2, x2**2, np.sin(x1), np.cos(x1), np.sin(x2), np.cos(x2),
                    np.sin^2(x1), np.cos^2(x1), np.sin^2(x2), np.cos^2(x2), np.exp(x1), np.exp(x2)]

        X = np.concatenate(columns, axis=1)
        return X

features = feature_expand(data)
print("Dataset size:", np.shape(data))
print("Expanded dataset size:", np.shape(features))
```

Dataset size: (130, 2)
Expanded dataset size: (130, 15)

Logistic Regression

Use SciKit-Learn's Logistic Regression model to learn the decision boundary for this data, using regularization. (The C argument controls regularization strength.)

Train this model on your expanded feature set.

Details about how to use this are here: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html
(https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html)

Notes:

- λ is related to sklearn's regularization strength C by: $\lambda = 1/C$

```
In [5]: from sklearn.linear_model import LogisticRegression

def get_logistic_regressor(features, classes, L = 1):
    # YOUR CODE GOES HERE
    # - Instantiate model with regularization
    # - Fit model to expanded data
    model = LogisticRegression(C = 1/L).fit(features, classes)
    return model
```

```
In [6]: for L in [1e-9, 1e-1, 1]:
        model = get_logistic_regressor(features, classes, L)
        plot_data(data, classes, **format, title=f"$\lambda={L}$")
        plot_contour(model.predict, feature_expand)
        plt.show()
```


d:\Python\lib\site-packages\sklearn\linear_model_logistic.py:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.


Increase the number of iterations (max_iter) or scale the data as shown in:

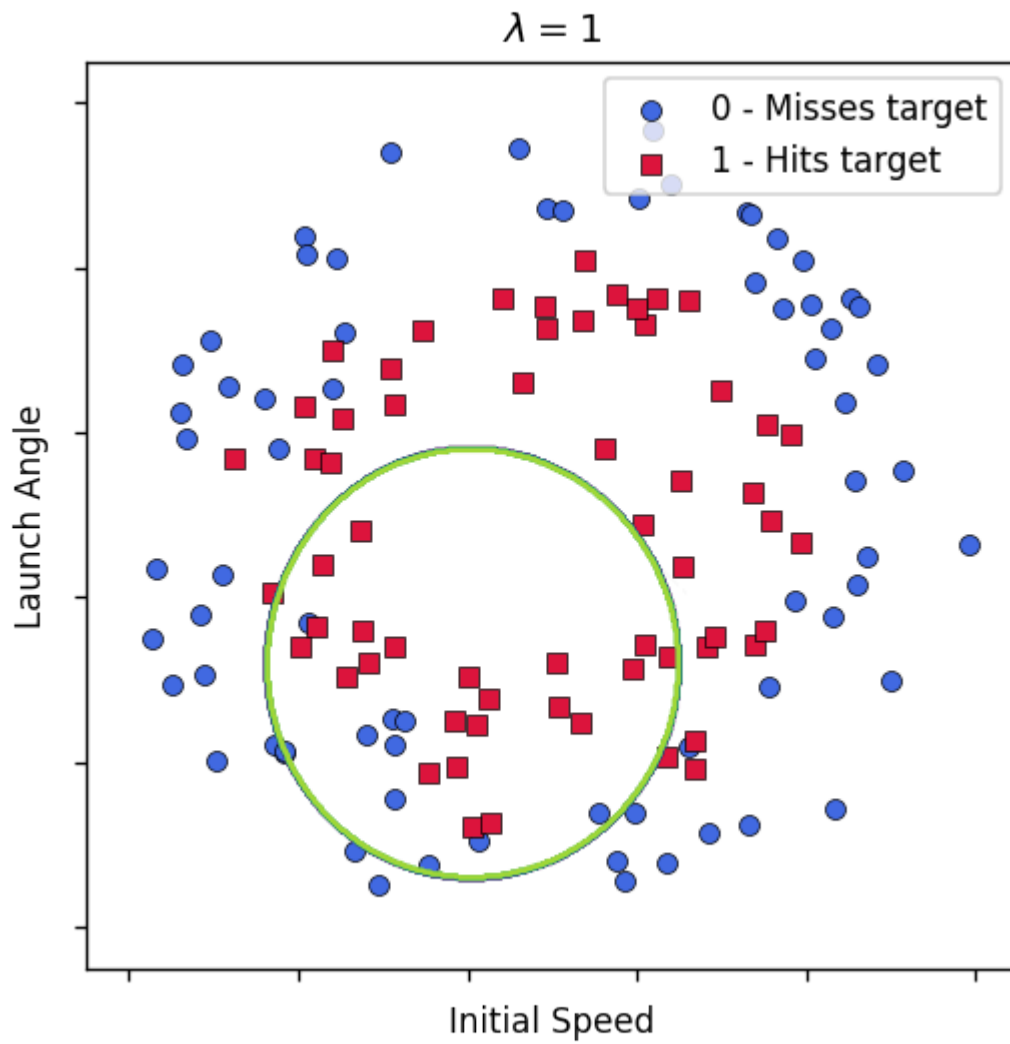
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
(https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,







As λ increases, note what happens to the decision boundary. Why does this occur?

As lambda increases, the classification becomes less accurate due to the decrease in regularization strength