

M8-L2 Problem 2

Let's revisit the material phase prediction problem once again. You will use this problem to try multi-class classification in PyTorch. You will have to write code for a classification network and for training.

In [201...

```
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
import torch
from torch import nn, optim

def plot_loss(train_loss, val_loss):
    plt.figure(figsize=(4,2),dpi=250)
    plt.plot(train_loss,label="Training")
    plt.plot(val_loss,label="Validation",linewidth=1)
    plt.legend()
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.show()

def split_data(X, Y):
    np.random.seed(100)
    N = len(Y)
    train_mask = np.zeros(N, dtype=np.bool_)
    train_mask[np.random.permutation(N)[:int(N*0.8)]] = True
    train_x, val_x = torch.Tensor(X[train_mask,:]), torch.Tensor(X[np.logical_not(train_mask)])
    train_y, val_y = torch.Tensor(Y[train_mask]), torch.Tensor(Y[np.logical_not(train_mask)])
    return train_x, val_x, train_y, val_y

x1 = np.array([7.4881350392732475,16.351893663724194,22.427633760716436,29.04883182
x2 = np.array([0.11120957227224215,0.1116933996874757,0.14437480785146242,0.1181820
X = np.vstack([x1,x2]).T
y = np.array([0,2,2,2,2,2,0,2,2,2,2,2,0,0,2,0,1,2,0,0,1,1,1,2,0,1,0,1,1,1,0,0,1,1,1

X = torch.Tensor(X)
Y = torch.tensor(y,dtype=torch.long)

train_x, val_x, train_y, val_y = split_data(X,Y)

def plot_data(newfig=True):
    xlim = [0,52.5]
    ylim = [0,1.05]
    markers = [dict(marker="o", color="royalblue"), dict(marker="s", color="crimson")]
    labels = ["Solid", "Liquid", "Vapor"]

    if newfig:
        plt.figure(figsize=(6,4),dpi=250)

    x = X.detach().numpy()
    y = Y.detach().numpy().flatten()
```

```

for i in range(1+max(y)):
    plt.scatter(x[y==i,0], x[y==i,1], s=40, **(markers[i]), edgecolor="black",


plt.scatter(val_x[:,0], val_x[:,1],s=120,c="None",marker="o",edgecolors="black")

plt.title("Phase of simulated material")
plt.legend(loc="upper right")
plt.xlim(xlim)
plt.ylim(ylim)
plt.xlabel("Temperature, K")
plt.ylabel("Pressure, atm")
plt.box(True)

def plot_model(model, res=200):
    xlim = [0,52.5]
    ylim = [0,1.05]
    xvals = np.linspace(*xlim,res)
    yvals = np.linspace(*ylim,res)
    x,y = np.meshgrid(xvals,yvals)
    XY = np.concatenate((x.reshape(-1,1),y.reshape(-1,1)),axis=1)
    XY = torch.Tensor(XY)
    color = model.predict(XY).reshape(res,res).detach().numpy()
    cmap = ListedColormap(["lightblue","lightcoral","palegreen"])
    plt.pcolor(x, y, color, shading="nearest", zorder=-1, cmap=cmap,vmin=0,vmax=2)
    return

plot_data()
plt.show()

```



Model definition

In the cell below, complete the definition for `PhaseNet`, a classification neural network.

- The network should take in 2 inputs and return 3 outputs.
- The network size and hidden layer activations are up to you.
- Make sure to use the proper activation function (for multi-class classification) at the final layer.
- The `predict()` method has been provided, to return the integer class value. You must finish `__init__()` and `forward()`.

In [202...

```
class PhaseNet(nn.Module):
    def __init__(self):
        super().__init__()
        # YOUR CODE GOES HERE
        self.lin1 = nn.Linear(2, 40)
        self.lin2 = nn.Linear(20, 40)
        self.lin3 = nn.Linear(40, 40)
        self.lin4 = nn.Linear(40, 20)
        self.lin5 = nn.Linear(40, 3)

        self.act1 = nn.ReLU()
        self.act2 = nn.Softmax()

    def predict(self,X):
        Y = self(X)
        return torch.argmax(Y,dim=1)

    def forward(self,X):
        # YOUR CODE GOES HERE
        X = self.lin1(X)
        X = self.act1(X)

        # X = self.lin2(X)
        # X = self.act1(X)

        X = self.lin3(X)
        X = self.act1(X)

        # X = self.lin4(X)
        # X = self.act1(X)

        X = self.lin5(X)
        X = self.act2(X)
        return X
```

Training

Most of the training code has been provided below. Please add the following where indicated:

- Define a loss function (for multiclass classification)
- Define an optimizer and call it `opt`. You may choose which optimizer.

Make sure the training curves you get are reasonable.

In [203...

```
model = PhaseNet()

lr = 0.001
epochs = 1000

# Define loss function
# YOUR CODE GOES HERE
lossfun = nn.CrossEntropyLoss()

# Define an optimizer, `opt`
# YOUR CODE GOES HERE
opt = optim.Adam(params=model.parameters(), lr=lr)

train_hist = []
val_hist = []


for epoch in range(epochs+1):
    model.train()
    loss_train = lossfun(model(train_x), train_y)
    train_hist.append(loss_train.item())

    model.eval()
    loss_val = lossfun(model(val_x), val_y)
    val_hist.append(loss_val.item())

    opt.zero_grad()
    loss_train.backward()
    opt.step()
    if epoch % int(epochs / 25) == 0:
        print(f"Epoch {epoch:>4} of {epochs}:   Train Loss = {loss_train.item():.4f}")

plot_loss(train_hist, val_hist)
```

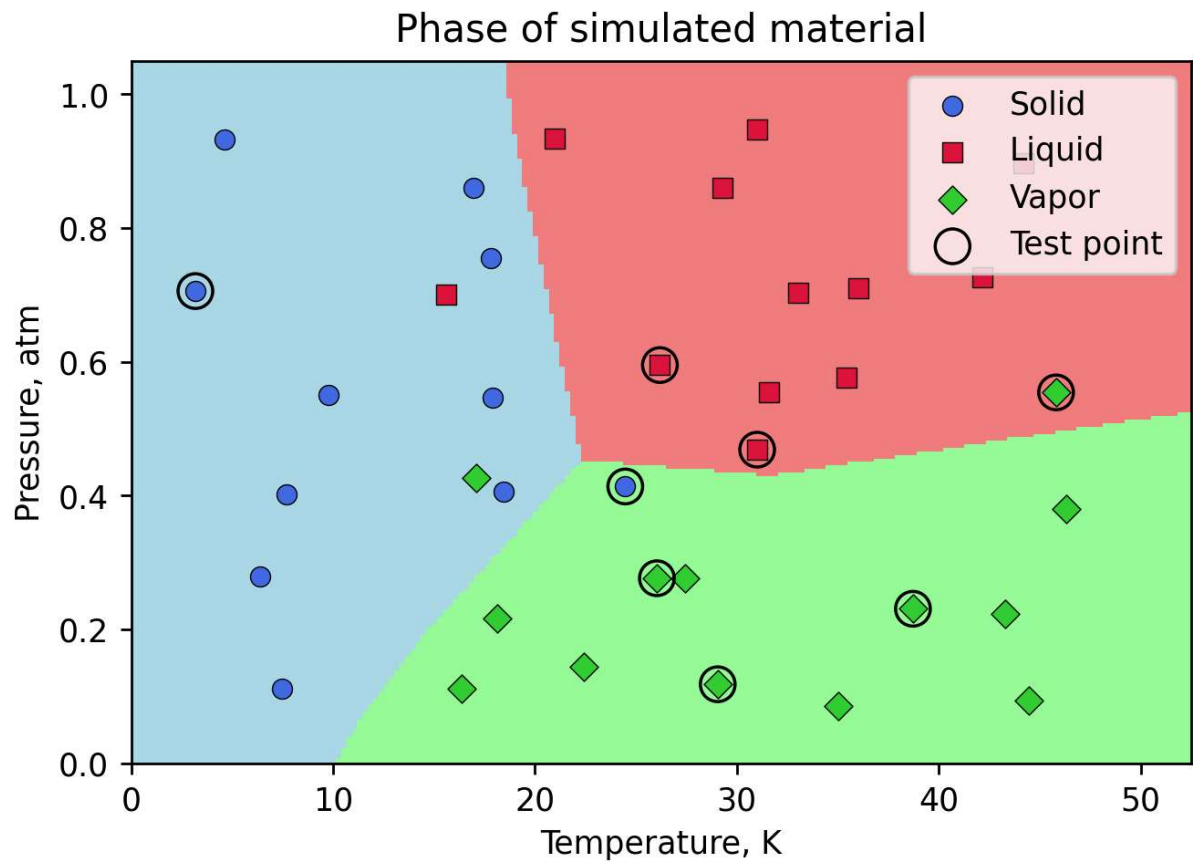
Epoch 0 of 1000:	Train Loss = 1.2509	Validation Loss = 1.3134
Epoch 40 of 1000:	Train Loss = 1.0531	Validation Loss = 1.0737
Epoch 80 of 1000:	Train Loss = 1.0194	Validation Loss = 1.0348
Epoch 120 of 1000:	Train Loss = 0.9556	Validation Loss = 0.9810
Epoch 160 of 1000:	Train Loss = 0.8808	Validation Loss = 0.9252
Epoch 200 of 1000:	Train Loss = 0.8180	Validation Loss = 0.8775
Epoch 240 of 1000:	Train Loss = 0.7596	Validation Loss = 0.8401
Epoch 280 of 1000:	Train Loss = 0.7200	Validation Loss = 0.8141
Epoch 320 of 1000:	Train Loss = 0.6969	Validation Loss = 0.7992
Epoch 360 of 1000:	Train Loss = 0.6824	Validation Loss = 0.7904
Epoch 400 of 1000:	Train Loss = 0.6723	Validation Loss = 0.7859
Epoch 440 of 1000:	Train Loss = 0.6648	Validation Loss = 0.7835
Epoch 480 of 1000:	Train Loss = 0.6590	Validation Loss = 0.7831
Epoch 520 of 1000:	Train Loss = 0.6543	Validation Loss = 0.7813
Epoch 560 of 1000:	Train Loss = 0.6504	Validation Loss = 0.7830
Epoch 600 of 1000:	Train Loss = 0.6471	Validation Loss = 0.7858
Epoch 640 of 1000:	Train Loss = 0.6442	Validation Loss = 0.7879
Epoch 680 of 1000:	Train Loss = 0.6418	Validation Loss = 0.7906
Epoch 720 of 1000:	Train Loss = 0.6397	Validation Loss = 0.7926
Epoch 760 of 1000:	Train Loss = 0.6378	Validation Loss = 0.7950
Epoch 800 of 1000:	Train Loss = 0.6362	Validation Loss = 0.7962
Epoch 840 of 1000:	Train Loss = 0.6348	Validation Loss = 0.7970
Epoch 880 of 1000:	Train Loss = 0.6333	Validation Loss = 0.7888
Epoch 920 of 1000:	Train Loss = 0.6323	Validation Loss = 0.7914
Epoch 960 of 1000:	Train Loss = 0.6313	Validation Loss = 0.7923
Epoch 1000 of 1000:	Train Loss = 0.6305	Validation Loss = 0.7924



Plot results

Plot your network predictions with the data by running the following cell. If your network has significant overfitting/underfitting, go back and retrain a new network with different layer sizes/activations.

```
In [204... plot_data(newfig=True)  
plot_model(model)  
plt.show()
```



```
In [ ]:
```